© 2008 Society for Industrial and Applied Mathematics

A DETERMINISTIC SUBEXPONENTIAL ALGORITHM FOR SOLVING PARITY GAMES*

MARCIN JURDZIŃSKI†, MIKE PATERSON†, AND URI ZWICK‡

Abstract. The existence of polynomial-time algorithms for the solution of parity games is a major open problem. The fastest known algorithms for the problem are *randomized* algorithms that run in subexponential time. These algorithms are all ultimately based on the randomized subexponential simplex algorithms of Kalai and of Matoušek, Sharir, and Welzl. Randomness seems to play an essential role in these algorithms. We use a completely different, and elementary, approach to obtain a *deterministic* subexponential algorithm for the solution of parity games. The new algorithm, like the existing randomized subexponential algorithms, uses only polynomial space, and it is almost as fast as the randomized subexponential algorithms mentioned above.

1. Introduction. A parity game [11, 15] is played on a directed graph (V, E) by two players, *Even* and *Odd*, who move a token from vertex to vertex along the edges of the graph so that an infinite path is formed. A partition (V_0, V_1) is given of the set V of vertices: player Even moves if the token is at a vertex of V_0, and player Odd moves if the token is at a vertex of V_1. Finally, a priority function $p : V \to \mathbb{N}$ is given. The players compete for the parity of the highest priority occurring infinitely often: player Even wins if $\limsup_{i \to \infty} p(v_i)$ is even, while player Odd wins if it is odd, where $v_0, v_1, v_2, \dots$ is the infinite path formed by the players.

The algorithmic problem of solving parity games is, given a parity game $G = (V_0, V_1, E, p)$ and an initial vertex $v_0 \in V$, to determine whether player Even has a winning strategy in the game if the token is initially placed on vertex v_0. Algorithms for solving parity games [33, 20, 32, 15, 2] usually compute the winning sets win_0 and win_1, i.e., the sets of vertices from which players Even and Odd, respectively, have a winning strategy. By the determinacy theorem for parity games [11, 15] the winning sets win_0 and win_1 form a partition of the set of vertices V. None of these algorithms is known to run in polynomial time and the existence of a polynomial-time algorithm for the solution of parity games is a long-standing open problem [12, 15].

The original motivation for the study of parity games comes from the area of formal verification of systems by temporal logic model checking [5, 15]. The problem of solving parity games is polynomial-time equivalent to the nonemptiness problem of ω-automata on infinite trees with Rabin-chain acceptance conditions [12], and to the model checking problem of the modal μ-calculus (modal fixpoint logic) which is a formalism of great expressiveness and succinctness in formal specification and

†Department of Computer Science, University of Warwick, Coventry CV4 7AL, UK (Marcin.Jurdzinski@dcs.warwick.ac.uk, msp@dcs.warwick.ac.uk).

‡School of Computer Science, Tel Aviv University, Tel Aviv 69978, Israel (zwick@cs.tau.ac.il).

validation [10, 15]. The model checking problem is a fundamental algorithmic problem in automated hardware and software verification [10, 5].

Another important motivation to study the problem of solving parity games is its intriguing complexity theoretic status: the problem is known to be in NP ∩ co-NP [12] and even in UP ∩ co-UP [19] but, as mentioned, despite considerable efforts of the community [12, 20, 32, 15, 2, 28] no polynomial-time algorithm has been found so far. (The complexity class UP, a.k.a. unambiguous NP, is defined to contain all problems that can be recognized by an unambiguous nondeterministic polynomial-time Turing machine. A Turing machine is unambiguous if for every input it has at most one accepting computation. Clearly, the inclusions P ⊆ UP ⊆ NP hold.) Moreover, parity games are polynomial-time reducible to mean payoff games [34], simple stochastic games [6], and discounted payoff games [6, 34]. A stochastic generalization of parity games was also studied [9, 3]. The problems of solving all those games are in UP ∩ co-UP as well [19, 3]. Condon has shown that simple stochastic games are complete (with respect to log-space reductions) in the class of log-space randomized alternating Turing machines [6].

The task of solving parity, mean payoff, discounted payoff, and simple stochastic games can be also viewed as a search problem: given a game graph, compute optimal strategies for both players. The value functions used in strategy improvement algorithms [7, 25, 32, 2] witness membership of all those optimal strategies search problems in PLS (i.e., the class of polynomial local search problems) [17]. On the other hand, the problem of computing optimal strategies in simple stochastic games can be reduced in polynomial time to solving a P-matrix linear complementarity problem [14, 31, 21], and to finding a Brouwer fixpoint [18], and hence it is also in PPAD (see [29] for a definition of PPAD). It follows that there are polynomial-time reductions from the problems of computing optimal strategies in parity, mean payoff, discounted payoff, and simple stochastic games to the problem of finding Nash equilibria in bimatrix games [8, 4].

Let $n = |V|$ and $m = |E|$ be the numbers of vertices and edges of a parity game graph and let d be the number of different priorities assigned to vertices by the priority function $p : V \to \mathbb{N}$. For parity games with a small number of priorities, more specifically if $d = O(n^{1/2})$, the progress-measure lifting algorithm [20] gave, until recently, the best time complexity of $O(dm(2n/d)^{d/2})$. This has been improved by Schewe [30] to $O(m(\kappa n/d)^{d/3})$, where $\kappa \le (2e)^{3/2}$. If $d = \Omega(n^{(1/2)+\varepsilon})$, then the randomized algorithm of Björklund, Sandberg, and Vorobyov [2] has a better (expected) running-time bound of $n^{O(\sqrt{n/\log n})}$.

The main contribution of this paper is a *deterministic* algorithm for solving parity games which achieves roughly the same complexity as the randomized algorithm of Björklund, Sandberg, and Vorobyov [2]: the complexity of our algorithm is $n^{O(\sqrt{n/\log n})}$ if the out-degree of all vertices is bounded, and is $n^{O(\sqrt{n})}$ otherwise. The new algorithm uses only polynomial space.

The randomized algorithm of Björklund, Sandberg, and Vorobyov [2] is based on the randomized algorithm of Ludwig [25] for simple stochastic games, which in turn is inspired by the subexponential randomized simplex algorithms for linear programming and LP-type problems by Kalai [23] and by Matoušek, Sharir, and Welzl [26]. For games with out-degree two these algorithms are instantiations of the Random-Facet algorithm for finding the unique sink in an acyclic unique sink orientation (AUSO) of a hypercube [13]. The nodes of a hypercube correspond to positional strategies for one of the players, and the orientation of an edge connecting two positional strategies

that differ at exactly one vertex is determined by which of the two strategies has a better value when the opponent plays a best-response strategy [7, 25, 32, 2].

In contrast, our deterministic algorithm for parity games is obtained by a modification of a more elementary algorithm of McNaughton and Zielonka for parity games [33, 15]. The methods we use are thus very different from those of Ludwig and Björklund et al. Our method is applicable, so it seems, only to parity games, while the randomized algorithms for finding the unique sink in an AUSO [13] and for solving an LP-type problem [26] can be applied to a number of problems, including computing the values of parity, mean payoff, discounted payoff, and simple stochastic games [2, 1, 16].

The recent improvement of the complexity of solving parity games with a small number of priorities due to Schewe [30] was inspired by the preliminary version of this paper published at SODA'06 [22]. Schewe refined our technique of searching and removing dominions while running the classical recursive algorithm [27, 33]: instead of a brute-force search he used a modification of the progress measure lifting algorithm [20].

2. Definitions. A *parity game* $G = (V_0, V_1, E, p)$ is composed of two disjoint sets of vertices V_0 and V_1, a set of directed edges $E \subseteq V \times V$, where $V = V_0 \cup V_1$, and a *priority* function $p : V_0 \cup V_1 \to \mathbb{N}$, defined on its vertices. Every vertex $u \in V$ has at least one outgoing edge $(u, v) \in E$. The game is played by two players: *Even*, also referred to as Player 0, and *Odd*, also referred to as Player 1.

The game starts at some vertex $v_0 \in V$. The players construct an infinite path (a play) as follows. Let u be the last vertex added so far to the path. If $u \in V_0$, then Player 0 chooses an edge $(u, v) \in E$. Otherwise, if $u \in V_1$, then Player 1 chooses an edge $(u, v) \in E$. In either case, vertex v is added to the path, and a new edge is then chosen by either Player 0 or Player 1. As each vertex has at least one outgoing edge, the path constructed can always be continued.

Let $v_0, v_1, v_2, \ldots$ be the infinite path constructed by the two players, and let $p(v_0), p(v_1), p(v_2), \ldots$ be the sequence of the priorities of the vertices on the path. Player 0 wins the play if the largest priority seen infinitely many times is even, and Player 1 wins otherwise. Observe that removing an arbitrary finite prefix of a play in a parity game does not change the winner; we refer to this property of parity games as *prefix independence*.

A *strategy* for Player i in a game G specifies, for every finite path $v_0, v_1, \ldots, v_k$ in G that ends in a vertex $v_k \in V_i$, an edge $(v_k, v_{k+1}) \in E$. A strategy is said to be a *positional* strategy if the edge $(v_k, v_{k+1}) \in E$ chosen depends only on v_k, the last vertex visited. A strategy for Player i is said to be a *winning* strategy if using this strategy ensures a win for Player i, no matter which strategy is used by the other player. The determinacy theorem for parity games [11, 15] says that for every parity game G and every start vertex v_0, either Player 0 has a winning strategy or Player 1 has a winning strategy. (This claim is not immediate as the games considered are infinite.) Furthermore, if a player has a winning strategy from a vertex in a parity game, then she also has a winning positional strategy from this vertex.

In the parity game G illustrated in Figure 2.1, the initial vertex is labeled a, *Even*'s vertices are represented as squares (even number of sides), and *Odd*'s as triangles. The numbers within the vertices show priorities. Note that each player's vertices may have both even and odd parities. As an example of a play in G, if each player were to choose the double-headed arrow out of each of their vertices, then the infinite path formed would be $a, d, b, e, d, b, e, \ldots$, and the largest priority seen infinitely often would be 4 at vertex e. So *Even* would win this play.

FIG. 2.1. *A parity game G. Double-headed arrows show a positional strategy for each player.*

The winning set for Player i, denoted by $win_i(G)$, is the set of vertices of the game from which Player i has a winning strategy. By the determinacy theorem for parity games [11, 15] we have that $win_0(G) \cup win_1(G) = V$.

3. Overview of the new algorithm. The previously known deterministic algorithm [33, 15] on which our improvement is built will be described fully in section 5. It has a recursive structure: solving a game with n vertices may require two recursive calls to smaller games. In the worst case, each of these games may have $n-1$ vertices, resulting in a running time satisfying the recurrence $T(n) \leq 2T(n-1) + O(n^2)$, which yields $T(n) = O(2^n)$. We offer no improvement in the first of the two recursive calls, but we do take advantage of a special feature of the second of these.

We introduce the notion of a *dominion*. An i-dominion, as the name suggests, is a set of vertices D *ruled* by Player i, in the sense that Player i can win from every vertex of D, without leaving D and without allowing the other player to leave D. One example of an i-dominion would be the whole of $win_i(G)$, but there may well be other "i-closed" subsets of $win_i(G)$ which are i-dominions. Although finding i-dominions can be just as hard as finding $win_i(G)$, we show that searching for small enough dominions is feasible, though taking time exponential in the size of dominion sought.

The significance of dominions for our algorithm depends on two properties. First, every i-dominion found can be easily removed at small cost, leaving a smaller game to be solved. Second, the second of the recursive calls is to a game resulting from the removal of a dominion. Therefore, if we look for and then remove all small dominions before entering the recursive calls, we can be sure that the second recursive call is to a substantially reduced game. The corresponding recurrence is then of the form $T(n) \leq T(n-1) + T(n-\ell) + O(n^\ell)$ for some $\ell = \ell(n)$.

With an appropriate choice of $\ell(n)$ to achieve a balance between the time to search for dominions of size up to ℓ and the savings from avoidance of the worst cases of the second recursive call, we achieve our subexponential algorithm with running time $n^{O(\sqrt{n})}$.

In the next section we prepare for the algorithms by introducing some key notions (*i-closed* and *reachability* set) and proving some of their properties. Lemmas 4.5 and 4.6 lay the foundation for the exponential algorithm. They show that we can begin to solve a game G by considering the set A of vertices with highest priority which Player i (say) would like the play to visit infinitely often, and identifying the set A^* from which Player i can guarantee to reach A at least once. From Lemma 4.6, we see that by first solving the smaller game G' based on vertices in $V(G) \setminus A^*$ we can identify a subset U of the winning set of the opponent of Player i, say Player j, in G.

By Lemma 4.5 we then know that the set U^*, from which Player j can guarantee to reach U at least once, is also included in the winning set of Player j in G. Moreover, Lemma 4.5 establishes that the winning set of Player i in G is equal to her winning set in the smaller game based on vertices in $V(G) \setminus U^*$, and hence the task of solving the game G is reduced to the task of solving the smaller game. For an illustration, see Figure 5.2, where $A^* = reach_i(A)$, $U = W_j'$, and $U^* = reach_j(W_j')$.

After giving details of the exponential algorithm in section 5, we show in section 6 how to find dominions. In sections 7 and 8 we integrate this search-and-remove process into our new algorithm and analyze the resulting running time.

4. Preliminaries. The results presented in this section are well known [15] and form the basis of algorithms by McNaughton [27] and Zielonka [33]. We include our detailed exposition of them here in order to fix the terminology and to make the paper self-contained.

A set $B \subseteq V$ is said to be *i-closed*, where $i \in \{0, 1\}$, if for every $u \in B$ the following hold:

- if $u \in V_i$, then there is some $(u, v) \in E$, such that $v \in B$; and
- if $u \in V_{\neg i}$, then for every $(u, v) \in E$, we have $v \in B$.

(We use $\neg i$ for the element $(1 - i)$ in $\{0, 1\}$.) In other words, a set B is i-closed if Player i can always choose to stay in B while Player $\neg i$ cannot escape from it, i.e., B is a "trap" for Player $\neg i$.

LEMMA 4.1. *For each $i \in \{0, 1\}$, the set $win_i(G)$ is i-closed.*

Proof. The proof is straightforward from the definitions and prefix independence of parity games. ☐

Let $A \subseteq V$ be an arbitrary set. The *i-reachability set* of A, denoted $reach_i(A)$, contains all vertices in A and all vertices from which Player i has a strategy to enter the set A at least once; we call such a strategy an *i-reachability strategy to* set A. (See Figure 4.1 for a simple example.)


FIG. 4.1. *The 0-reachability set of A and a positional 0-reachability strategy to set A.*

LEMMA 4.2. *For every set $A \subseteq V$ and $i \in \{0, 1\}$, the set $V \setminus reach_i(A)$ is $(\neg i)$-closed.*

Proof. Let $u \in V \setminus reach_i(A)$. Recall that every vertex has at least one outgoing edge; hence if $u \in V_{\neg i}$, then there must be an edge $(u, v) \in E$ from vertex u into the set $V \setminus reach_i(A)$, i.e., such that $v \notin reach_i(A)$, since otherwise vertex u would be in $reach_i(A)$. Similarly, if $u \in V_i$, then all edges from vertex u must go into the set $V \setminus reach_i(A)$. Therefore, the set $V \setminus reach_i(A)$ is $(\neg i)$-closed. ☐

LEMMA 4.3. *For every set $A \subseteq V$ and $i \in \{0, 1\}$, the set $reach_i(A)$ can be computed in $O(m)$ time, where $m = |E|$ is the number of edges in the game.*

Proof. The vertices of A are in $reach_i(A)$ so we initialize $B \leftarrow A$. We then iteratively add to B every vertex of V_i that has at least one edge going into B, and every vertex of $V_{\neg i}$ all of whose edges go into B. We stop when no new vertices can be added to B.

Some care is needed to keep the time in $O(m)$. One method is to maintain an adjacency list giving, for each vertex, the *incoming* edges together with a count of the number of outgoing edges. At each step we take an edge (u, v) entering B and delete it (from the list of edges entering v and from the count of edges leaving u): if $u \in B$, then nothing more is done; otherwise, if $u \in V_i$, then u is added to B; otherwise (when $u \in V_{\neg i} \setminus B$), if there are no other edges from u, then u is added to B.

It is easy to see that this process can be performed in $O(m)$ time, and that when it ends we have $B = reach_i(A)$, as required.   ☐

If $B \subseteq V$ is such that for every vertex $u \in V \setminus B$ there is an edge (u, v) with $v \in V \setminus B$, then the *subgame* $G \setminus B$ is the game obtained from G by removing the vertices of B and all the edges that touch them. We will only be using B's for which $V \setminus B$ is an i-closed set for some i. In such cases $G \setminus B$ is always well-defined. The next lemmas show some useful properties of subgames.

LEMMA 4.4. *Let G' be a subgame of G and let $i \in \{0, 1\}$. If V', the vertex set of G', is i-closed in G, then $win_i(G') \subseteq win_i(G)$.*

Proof. A winning strategy for Player i from the set $win_i(G')$ in the subgame G' is also winning for her from the same set in the original game G. Player $\neg i$ cannot escape to $V \setminus V'$, since the set V' is i-closed in G.   ☐

The following lemma implies that if we know an arbitrary nonempty subset U of the winning set of a player, say Player j, in a game G, then computing the winning sets of both players in G can be reduced to computing their winning sets in the smaller game $G \setminus reach_j(U)$.

LEMMA 4.5. *Let G be a parity game, let $i \in \{0, 1\}$ and $j = \neg i$. If $U \subseteq win_j(G)$ and $U^* = reach_j(U)$, then $win_j(G) = U^* \cup win_j(G \setminus U^*)$ and $win_i(G) = win_i(G \setminus U^*)$.*

FIG. 4.2. *Diagram illustrating Lemma* 4.5.

Proof. Let $W_j = U^* \cup win_j(G \setminus U^*)$ and $W_i = win_i(G \setminus U^*)$; see Figure 4.2. Since (W_i, W_j) is a partition of V, it suffices to show that $W_i \subseteq win_i(G)$ and $W_j \subseteq win_j(G)$. By Lemma 4.2, $V \setminus U^*$, the vertex set of $G \setminus U^*$, is i-closed. The first inclusion then follows from Lemma 4.4.

To show the second inclusion, we exhibit a strategy for Player j that is winning for her from the set W_j in the game G. By the assumption that $U \subseteq win_j(G)$, there is a strategy σ for Player j in the game G which is winning for her from all vertices in U. Let τ be a winning strategy for Player j from the set $win_j(G \setminus U^*)$ in the

subgame $G \setminus U^*$. A strategy π for Player j in the game G is made by composing strategies τ and σ in the following way: if the play so far is contained in the set $win_j(G \setminus U^*)$, then follow strategy τ; otherwise use the j-reachability strategy to the set U and "restart" the play following the strategy σ thenceforth. The strategy π is well-defined because, by Lemma 4.1, Player i can escape from $win_j(G \setminus U^*)$ only into the set U^*. By prefix independence of parity games, the strategy π is a winning strategy for Player j, because if it ever switches from following τ to following σ, then an infinite suffix of the play is winning for Player j.   ☐

The next lemma complements Lemma 4.5 by providing an algorithmic method which either finds a nonempty subset of the winning set of a player, say Player j, in a parity game G, or (if it returns an empty set) concludes that Player $\neg j$ can win from every vertex in G.

LEMMA 4.6. *Let G be a parity game. Let $d = d(G)$ be the highest priority and let $A = A_d(G)$ be the set of vertices of highest priority. Let $i = d \bmod 2$ and $j = \neg i$. Let $G' = G \setminus reach_i(A)$. Then we have $win_j(G') \subseteq win_j(G)$. Also, if $win_j(G') = \emptyset$, then $win_i(G) = V(G)$, i.e., Player i wins from every vertex of G.*

(As an example, consider Figures 4.1 and 4.3, with $i = 0$ and $j = 1$.)

FIG. 4.3. *The game $G' = G \setminus reach_0(A)$ and winning sets $W'_i = win_i(G')$ for $i = 0, 1$.*

Proof. That $win_j(G') \subseteq win_j(G)$ follows from Lemmas 4.2 and 4.4.

Suppose now that $win_j(G') = \emptyset$. Let τ be a winning strategy for Player i from $win_i(G')$ (which, by determinacy, is equal to $V \setminus reach_i(A)$) in the subgame G'. We construct a strategy π for Player i in the following way: if a play so far is contained in the set $win_i(G')$, then follow strategy τ; otherwise the current vertex is in $reach_i(A)$ so follow the i-reachability strategy to the set A; moreover, each time the play reenters the set $win_i(G')$ "restart" the play and follow strategy τ, etc. If a play following the strategy π visits $reach_i(A)$ (and hence A) infinitely often, then it is winning for Player i because $i = d \bmod 2$. Otherwise, it has an infinite suffix played according to strategy τ, and hence it is winning for Player i by prefix independence of parity games.   ☐

5. An exponential algorithm. A simple exponential-time algorithm for the solution of parity games is given in Figure 5.1. This algorithm originates from the work of McNaughton [27] and was first presented for parity games by Zielonka [33, 15]. Algorithm **win**(G) receives a parity game G and returns the pair of winning sets $(win_0(G), win_1(G))$ for the two players.

Algorithm **win**(G) is based on Lemmas 4.5 and 4.6. It starts by letting d be the largest priority in G and by letting A be the set of vertices having this highest priority. Let $i = d \bmod 2$ be the index of the player associated with the highest priority, and

```
algorithm win(G)

if V(G) = ∅ then return (∅, ∅)

d ← d(G) ; A ← A_d(G)
i ← d mod 2 ; j ← ¬i

(W'_0, W'_1) ← win(G \ reach_i(A))

if W'_j = ∅ then
    (W_i, W_j) ← (V(G), ∅)
else
    (W''_0, W''_1) ← win(G \ reach_j(W'_j))
    (W_i, W_j) ← (W''_i, V(G) \ W''_i)
endif

return (W_0, W_1)
```

FIG. 5.1. *An exponential algorithm for solving parity games.*

FIG. 5.2. *A game G and its subgames $G' = G \setminus reach_i(A)$ and $G'' = G \setminus reach_j(W'_j)$.*

let $j = \neg i$ be the index of the other player. The algorithm first finds the winning sets (W'_0, W'_1) of the smaller game $G' = G \setminus reach_i(A)$, using a recursive call; see Figure 5.2. By Lemma 4.6, if $W'_j = \emptyset$, then Player i wins from all vertices of G and we are done. Otherwise, again by Lemma 4.6, we know that $W'_j \subseteq win_j(G)$. The algorithm then finds the winning sets (W''_0, W''_1) of the smaller game $G'' = G \setminus reach_j(W'_j)$ by a second recursive call. By Lemma 4.5, we then know that $win_i(G) = W''_i$ and $win_j(G) = reach_j(W'_j) \cup W''_j = V(G) \setminus W''_i$.

A small detailed illustration of the main steps of the algorithm is given in Figures 4.1, 4.3, 5.3, and 5.4.

THEOREM 5.1. *Algorithm* win(G) *correctly finds the winning sets of the parity game G. Its running time is $O(2^n)$, where $n = |V(G)|$ is the number of vertices in G.*

Proof. The correctness of the algorithm follows from Lemmas 4.5 and 4.6, as argued above. Let $T(n)$ be the maximum running time of algorithm win(G) for a game on at most n vertices. Algorithm win(G) makes two recursive calls win(G') and win(G'') on games with at most $n-1$ vertices. Other than that, it performs only $O(n^2)$ operations. (The most time-consuming operations are the computations of the sets $reach_i(A)$ and $reach_j(W'_j)$.) Thus $T(n) \leq 2T(n-1) + O(n^2)$. It is easy to see then that $T(n) = O(2^n)$. □

FIG. 5.3. *The 1-reachability set of W_1'.*

FIG. 5.4. *The game $G'' = G \setminus \mathrm{reach}_1(W_1')$.*

6. Finding small dominions. A set $D \subseteq V(G)$ is said to be an *i-dominion* if Player i can win from every vertex of D without ever leaving D. Note, in particular, that an *i*-dominion must be *i*-closed. A set $D \subseteq V(G)$ is said to be a *dominion* if it is either a 0-dominion or a 1-dominion. By prefix independence of parity games, the winning set $win_i(G)$ of Player i is an *i*-dominion.

LEMMA 6.1. *Let G be a parity game on n vertices and let $\ell \leq n/3$. There is an $O(n^\ell)$-time algorithm that finds a nonempty dominion in G of size at most ℓ, or determines that no such dominion exists.*

Proof. If $\ell \leq n/3$, then, for all $j \leq \ell$, we have that $\binom{n}{j}/\binom{n}{j-1} > 2$. The number $\sum_{j=1}^{\ell} \binom{n}{j}$ of subsets of V of size at most ℓ is therefore at most $2\binom{n}{\ell}$. For each such subset U we check, in $O(\ell^2)$ time, whether it is 0-closed or 1-closed. If both tests fail, then U is clearly not a dominion. If U is *i*-closed, for some $i \in \{0,1\}$, we form the game $G[U]$, which is the game G restricted to U. This is well-defined since U is *i*-closed. We now apply the exponential algorithm of the previous section to $G[U]$ and find out, in $O(2^\ell)$ time, whether Player i can win from all the vertices of $G[U]$. If so, then U is an *i*-dominion; otherwise it is not. The total running time of the algorithm is therefore $O(\binom{n}{\ell}2^\ell) = O(n^\ell)$, as required. $\qquad \blacksquare$

In a game with bounded out-degrees we can find small dominions even faster. For simplicity, the lemma below and the analysis in section 8 are stated for games in which the out-degree of every vertex is exactly two. Note, however, that for every constant b, every game on n vertices with out-degrees at most b can be easily converted into an equivalent game on at most $n(b-1)$ vertices with out-degrees exactly two, by replacing each higher-degree vertex by a binary tree.

LEMMA 6.2. *Let G be a parity game on n vertices in which the out-degree of each vertex is two. There is an $O(n2^\ell \ell \log \ell)$-time algorithm that finds a nonempty dominion in G of size at most ℓ, or determines that no such dominion exists.*

Proof. Assume, without loss of generality, that the vertices of G are numbered from 1 to n. Let $u \in V$ be a vertex of G and let (u, v_0) and (u, v_1) be the two edges emanating from v, where $v_0 \leq v_1$. We say that (u, v_0) is the 0th outgoing edge of v, while (u, v_1) is the 1st outgoing edge of v.

The algorithm generates at most $O(n2^\ell)$ 0-closed sets of size at most ℓ that are candidates for being 0-dominions. For every vertex $v \in V$ and a binary sequence $\langle a_1, \ldots, a_\ell \rangle \in \{0,1\}^\ell$, construct a set $U \subset V$ as follows. Start with $U = \{v\}$ and $r = 1$. Vertices added to U are initially unmarked. As long as there is still an unmarked vertex in U, pick the smallest such vertex $u \in U$ and mark it. If $u \in V_0$, then add the endpoint of the a_rth outgoing edge of u to U, if it is not already there,

and increment r. If $u \in V_1$, then add the endpoints of both the outgoing edges of u to U. If at some stage $|U| > \ell$, then discard the set U and restart the construction with the next binary sequence.

If the process above ends with $|U| \le \ell$, then a 0-closed set of size at most ℓ has been found. Furthermore, for every vertex $u \in U \cap V_0$, one of the outgoing edges of u was selected. This corresponds to a suggested strategy for Player 0 in the game $G[U]$. Our algorithm therefore considers by exhaustive search all 0-closed sets of at most ℓ vertices, and for each set considers all possible positional strategies for Player 0.

Using an algorithm of King, Kupferman, and Vardi [24] we can check, in $O(\ell \log \ell)$ time, whether a given U and proposed strategy is indeed a winning strategy for Player 0 from all the vertices of U. Thus, if there is a 0-dominion of size at most ℓ in G, then the algorithm will find one. Finding 1-dominions of size at most ℓ can be done in an analogous manner. ☐

The algorithm described in Lemma 6.1 finds some i-dominion D if there is a dominion of size at most ℓ. We denote this algorithm by $\mathbf{dominion}(G, \ell)$, and suppose that it returns either the pair (D, i) if successful, or $(\emptyset, -1)$ if not.

7. The new subexponential algorithm. The new algorithm for solving parity games is given in Figure 7.1. The algorithm **new-win** starts by trying to find a dominion of size at most ℓ, where $\ell = \lceil \sqrt{2n} \rceil$ (and $\ell = \lceil \sqrt{n \log n} \rceil$ for games with bounded out-degree) is a parameter chosen to minimize the running time of the whole algorithm. If such a small i-dominion is found, then it is easy to remove it, as well as its i-reachability set, from the game and recurse on what is left over. If no small

```
algorithm new-win(G)

if V(G) = ∅ then return (∅, ∅)

n ← |V(G)| ; ℓ ← ⌈√2n⌉
(D, i) ← dominion(G, ℓ); j ← ¬i

if D = ∅ then
    (W₀, W₁) ← old-win(G)
else
    (W₀′, W₁′) ← new-win(G \ reachᵢ(D))
    (Wⱼ, Wᵢ) ← (Wⱼ′, V(G) \ Wⱼ′)
endif

return(W₀, W₁)
```

```
algorithm old-win(G)

d ← d(G) ; A ← A_d(G)
i ← d mod 2 ; j ← ¬i

(W₀′, W₁′) ← new-win(G \ reachᵢ(A))

if Wⱼ′ = ∅ then
    (Wᵢ, Wⱼ) ← (V(G), ∅)
else
    (W₀″, W₁″) ← new-win(G \ reachⱼ(Wⱼ′))
    (Wᵢ, Wⱼ) ← (Wᵢ″, V(G) \ Wᵢ″)
endif

return (W₀, W₁)
```

FIG. 7.1. *The new subexponential algorithm for solving parity games.*

dominion is found, then **new-win**(G) simply calls algorithm **old-win**(G), which is almost identical to the exponential algorithm **win**(G) of section 5. The only difference between **old-win**(G) and **win**(G) is that the recursive calls are made to **new-win**(G) and not to **win**(G).

THEOREM 7.1. *Algorithm* **new-win**(G) *correctly finds the winning sets of a parity game* G. *Its running time on a game with n vertices is* $n^{O(\sqrt{n})}$.

Proof. The correctness of the algorithm is immediate. We next analyze its running time. Let $T(n)$ be the maximum running time of **new-win**(G) on a game with at most n vertices.

Algorithm **new-win**(G) tries to find dominions of size at most $\ell = \lceil \sqrt{2n} \rceil$. By Lemma 6.1 this takes $O(n^\ell)$ time. If a nonempty dominion is found, then the algorithm simply proceeds on the remaining game, which has at most $n - 1$ vertices, and the remaining running time is therefore at most $T(n-1)$. Otherwise, a call to **old-win**(G) is made. This results in a call to **new-win**$(G \setminus reach_i(A))$, which takes at most $T(n-1)$ time. If the set W_j' returned by the call is empty, then we are done. Otherwise, $W_j' = win_j(G \setminus reach_i(A))$, and this is equal to $win_j(G)$ by Lemma 4.5. Therefore W_j' is a j-dominion of G. We are in the case that there is no small dominion in G, so we know that $|W_j'| > \ell$, and therefore the second recursive call **new-win**$(G \setminus reach_j(W_j'))$ takes at most $T(n - \ell)$ time. Thus we get

$$T(n) \leq O(n^\ell) + T(n-1) + T(n-\ell).$$

This recurrence relation, with $\ell = \lceil \sqrt{2n} \rceil$, is analyzed in Theorem 8.1, where it is shown that $T(n) = n^{O(\sqrt{n})}$. ☐

A slightly better bound is achieved for graphs with out-degree two.

THEOREM 7.2. *Consider the algorithm* **new-win**(G) *in which the variable ℓ is set to $\lceil \sqrt{n \log n} \rceil$. If the game G has n vertices and the out-degree of each of them is two, then the running time of the modified algorithm is* $n^{O(\sqrt{n/\log n})}$.

Proof. Note that if $\ell = \lceil \sqrt{n \log n} \rceil$, then $O(n 2^\ell \ell \log \ell) = n^{O(\sqrt{n/\log n})}$. Therefore, by Lemma 6.2 and by the analysis in the proof of the previous theorem, the time complexity $T(n)$ satisfies the following recurrence:

$$T(n) \leq n^{O(\sqrt{n/\log n})} + T(n-1) + T(n-\ell).$$

The recurrence is analyzed in the following section in Theorem 8.2, where we show that $T(n) = n^{O(\sqrt{n/\log n})}$. ☐

8. Solving the recurrence relations. In this section we analyze the recurrence relations used in the previous section to bound the running time of the new algorithm.

We start by analyzing the recurrence relation used to bound the running time of the algorithm for game graphs with arbitrary out-degrees.

THEOREM 8.1. *If $T(n)$ is a positive function such that, for every $n > 3$,*

$$T(n) \leq O(n^\ell) + T(n-1) + T(n-\ell),$$

where $\ell = \lceil \sqrt{2n} \rceil$, then $T(n) = n^{O(\sqrt{n})}$.

Proof. For every integer n we construct a binary tree $\mathcal{T}_n$ in the following way. The root of $\mathcal{T}_n$ is labeled by n. A node labeled by a number $k > 3$ has two children: a left child labeled by $k - 1$ and a right child labeled by $k - \lceil \sqrt{2k} \rceil$. Nodes labeled by the numbers 1, 2, and 3 are leaves. A node labeled by k has a cost of $k^{O(\sqrt{2k})}$

associated with it. It is easy to see that the sum of the costs of the nodes of $\mathcal{T}_n$ is an upper bound on $T(n)$.

Clearly, the length of every path in $\mathcal{T}_n$ from the root to a leaf is at most n. We say that such a path makes a *right turn* when it descends from a vertex to its right child. We next claim that each such path makes at most $\lfloor\sqrt{2n}\rfloor$ right turns. This follows immediately from the observation that the function $f(n) = n - \lceil\sqrt{2n}\rceil$ can be iterated on n at most $\lfloor\sqrt{2n}\rfloor$ times before reaching the value of 3 or less. This observation can be proved by induction, based on the fact that if $\frac{1}{2}j^2 < n \le \frac{1}{2}(j+1)^2$, then $n - \lceil\sqrt{2n}\rceil \le \frac{1}{2}j^2$. (Initially we have $j = \lfloor\sqrt{2n}\rfloor$ and finally, with $1 \le n \le 3$, we have $j \ge 1$.)

As each leaf of $\mathcal{T}_n$ is determined by the positions of the right turns on the path leading to it from the root, we get that the number of leaves in $\mathcal{T}_n$ is at most $\binom{n}{\lfloor\sqrt{2n}\rfloor}$. The total number of nodes in $\mathcal{T}_n$ is therefore at most $2\binom{n}{\lfloor\sqrt{2n}\rfloor}$. As the cost of each node is at most $n^{O(\sqrt{2n})}$, we immediately get that

$$T(n) \le 2\binom{n}{\lfloor\sqrt{2n}\rfloor}n^{O(\sqrt{2n})} = n^{O(\sqrt{n})},$$

as claimed. $\square$

A more careful analysis, in which the $O(n^\ell)$ term in the recurrence relation is replaced by $O\left(\binom{n}{\ell}2^\ell\right)$, can be used to show that $T(n) = O((cn)^{\sqrt{n/2}})$, for some constant $c > 0$, and that the choice $\ell = \lceil\sqrt{2n}\rceil$ is essentially optimal.

The running time of the algorithm for graphs with out-degree two satisfies a tighter recurrence relation, which is analyzed similarly in the next theorem.

THEOREM 8.2. *If $T(n)$ is a positive function such that, for every $n > 3$,*

$$T(n) \le n^{O(\sqrt{n/\log n})} + T(n-1) + T(n-\ell),$$

where $\ell = \lceil\sqrt{n\log n}\rceil$, then $T(n) = n^{O(\sqrt{n/\log n})}$.

Proof. The proof is similar to the proof of Theorem 8.1. For every integer n we again construct a tree $\mathcal{T}_n$. A node labeled by a number $k > 2$ now has a left child labeled by $k-1$ and a right child labeled by $k - \lceil\sqrt{k\log k}\rceil$. The cost of a node labeled by k is now $k^{O(k/\log k)}$. Every root-to-leaf path in $\mathcal{T}_n$ is again of length at most n, and it can now make at most $O(\sqrt{n/\log n})$ right turns. Thus, the number of nodes in $\mathcal{T}_n$ is at most $n^{O(\sqrt{n/\log n})}$. As the cost of each node is also $n^{O(\sqrt{n/\log n})}$, we get that $T(n) = n^{O(\sqrt{n/\log n})}$, as claimed. $\square$

9. Concluding remarks. We have obtained the first deterministic subexponential algorithm for solving parity games. Our algorithm does not seem to extend in an obvious way to the solution of the more general mean payoff games and simple stochastic games. On the other hand, the techniques that we have introduced in this paper have recently inspired a notable improvement in the running-time complexity of parity games with a small number of priorities [30].

REFERENCES

[1] H. BJÖRKLUND AND S. VOROBYOV, *A combinatorial strongly subexponential strategy improvement algorithm for mean payoff games*, Discrete Appl. Math., 155 (2007), pp. 210–229.

[2] H. BJÖRKLUND, S. SANDBERG, AND S. VOROBYOV, *A discrete subexponential algorithm for parity games*, in Proceedings of the 20th Annual Symposium on Theoretical Aspects of Computer Science (STACS), Lecture Notes in Comput. Sci. 2607, Springer, Berlin, 2003, pp. 663–674.

[3] K. CHATTERJEE, M. JURDZIŃSKI, AND T. A. HENZINGER, *Quantitative stochastic parity games*, in Proceedings of the 15th Annual ACM-SIAM Symposium on Discrete Algorithms (SODA), ACM, New York, SIAM, Philadelphia, 2004, pp. 114–123.

[4] X. CHEN AND X. DENG, *Settling the complexity of two-player Nash equilibrium*, in Proceedings of the 47th Annual Symposium on Foundations of Computer Science (FOCS), IEEE Computer Society Press, 2006, pp. 261–272.

[5] E. M. CLARKE, O. GRUMBERG, AND D. PELED, *Model Checking*, MIT Press, Cambridge, MA, 1999.

[6] A. CONDON, *The complexity of stochastic games*, Inform. and Comput., 96 (1992), pp. 203–224.

[7] A. CONDON, *On algorithms for simple stochastic games*, in Advances in Computational Complexity Theory, AMS, Providence, RI, 1993, pp. 51–73.

[8] C. DASKALAKIS, P. W. GOLDBERG, AND C. H. PAPADIMITRIOU, *The complexity of computing a Nash equilibrium*, in Proceedings of the 38th ACM Symposium on Theory of Computing (STOC), ACM Press, New York, 2006, pp. 71–78.

[9] L. DE ALFARO AND R. MAJUMDAR, *Quantitative solution of omega-regular games*, J. Comput. System Sci., 68 (2004), pp. 374–397.

[10] E. A. EMERSON, *Model checking and mu-calculus*, in Descriptive Complexity and Finite Models, N. Immerman and P. G. Kolaitis, eds., DIMACS Ser. Discrete Math. Theoret. Comput. Sci. 31, AMS, Providence, RI, 1996, pp. 185–214.

[11] E. A. EMERSON AND C. JUTLA, *Tree automata, μ-calculus, and determinacy*, in Proceedings of the 32nd Annual Symposium on Foundations of Computer Science (FOCS), IEEE Computer Society Press, 1991, pp. 368–377.

[12] E. A. EMERSON, C. S. JUTLA, AND A. P. SISTLA, *On model-checking for fragments of μ-calculus*, in Computer-Aided Verification (CAV), Lecture Notes in Comput. Sci. 697, Springer, Berlin, 1993, pp. 385–396.

[13] B. GÄRTNER, *The Random-Facet simplex algorithm on combinatorial cubes*, Random Structures Algorithms, 20 (2002), pp. 353–381.

[14] B. GÄRTNER AND L. RÜST, *Simple stochastic games and P-matrix generalized linear complementarity problems*, in Fundamentals of Computation Theory (FCT), Lecture Notes in Comput. Sci. 3623, Springer, Berlin, 2005, pp. 209–220.

[15] E. GRÄDEL, W. THOMAS, AND T. WILKE, EDS., *Automata, Logics, and Infinite Games. A Guide to Current Research*, Lecture Notes in Comput. Sci. 2500, Springer, Berlin, 2002.

[16] N. HALMAN, *Discrete and Lexicographic Helly Theorems and Their Relations to LP-type Problems*, Ph.D. thesis, Tel Aviv University, Tel Aviv, Israel, 2004.

[17] D. S. JOHNSON, C. H. PAPADIMITRIOU, AND M. YANNAKAKIS, *How easy is local search?*, J. Comput. System Sci., 37 (1988), pp. 79–100.

[18] B. JUBA, *On the Hardness of Simple Stochastic Games*, M.S. thesis, Carnegie Mellon University, Pittsburgh, PA, 2005.

[19] M. JURDZIŃSKI, *Deciding the winner in parity games is in UP ∩ co-UP*, Inform. Process. Lett., 68 (1998), pp. 119–124.

[20] M. JURDZIŃSKI, *Small progress measures for solving parity games*, in Proceedings of the 17th Annual Symposium on Theoretical Aspects of Computer Science (STACS), Lecture Notes in Comput. Sci. 1770, Springer, Berlin, 2000, pp. 290–301.

[21] M. JURDZIŃSKI AND R. SAVANI, *A simple P-matrix linear complementarity problem for discounted games*, in Computability in Europe (CiE), Lecture Notes in Comput. Sci. 5028, Springer, Berlin, 2008, pp. 283–293.

[22] M. JURDZIŃSKI, M. PATERSON, AND U. ZWICK, *A deterministic subexponential algorithm for solving parity games*, in Proceedings of the 17th Annual ACM-SIAM Symposium on Discrete Algorithms (SODA), ACM, New York, SIAM, Philadelphia, 2006, pp. 117–123.

[23] G. KALAI, *A subexponential randomized simplex algorithm (extended abstract)*, in Proceedings of the 24th Annual ACM Symposium on Theory of Computing (STOC), ACM Press, New York, 1992, pp. 475–482.

[24] V. KING, O. KUPFERMAN, AND M. Y. VARDI, *On the complexity of parity word automata*, in Foundations of Software Science and Computation Structures (FoSSaCS), Lecture Notes

in Comput. Sci. 2030, Springer, Berlin, 2001, pp. 276–286.

[25] W. Ludwig, *A subexponential randomized algorithm for the simple stochastic game problem*, Inform. and Comput., 117 (1995), pp. 151–155.

[26] J. Matoušek, M. Sharir, and E. Welzl, *A subexponential bound for linear programming*, Algorithmica, 16 (1996), pp. 498–516.

[27] R. McNaughton, *Infinite games played on finite graphs*, Ann. Pure Appl. Logic, 65 (1993), pp. 149–184.

[28] J. Obdržálek, *Fast mu-calculus model checking when tree-width is bounded*, in Computer-Aided Verification (CAV), Lecture Notes in Comput. Sci. 2725, Springer, Berlin, 2003, pp. 80–92.

[29] C. H. Papadimitriou, *On the complexity of the parity argument and other inefficient proofs of existence*, J. Comput. System Sci., 48 (1994), pp. 498–532.

[30] S. Schewe, *Solving parity games in big steps*, in Foundations of Software Technology and Theoretical Computer Science (FSTTCS), Lecture Notes in Comput. Sci. 4855, Springer, Berlin, 2007, pp. 449–460.

[31] O. Svensson and S. Vorobyov, *Linear complementarity and P-matrices for stochastic games*, in Perspectives of Systems Informatics, Andrei Ershov Memorial Conference (PSI 2006), Lecture Notes in Comput. Sci. 4378, Springer, Berlin, 2007, pp. 409–423.

[32] J. Vöge and M. Jurdziński, *A discrete strategy improvement algorithm for solving parity games (extended abstract)*, in Computer-Aided Verification (CAV), Lecture Notes in Comput. Sci. 1855, Springer, Berlin, 2000, pp. 202–215.

[33] W. Zielonka, *Infinite games on finitely coloured graphs with applications to automata on infinite trees*, Theoret. Comput. Sci., 200 (1998), pp. 135–183.

[34] U. Zwick and M. Paterson, *The complexity of mean payoff games on graphs*, Theoret. Comput. Sci., 158 (1996), pp. 343–359.